Chapter 6

Moral principles and moral theories

In Chapter 1, we defined a principle as a general rule or recommendation which applies to a number of specific cases. We also observed that principles can be closely related to moral concepts, a point which was reinforced in the last chapter when we considered whether, for example, justifying a moral claim on the basis of the concept of a right to life is just another way of saying that in most circumstances killing is wrong. This illustrates one way in which concepts and principles can be related – that they can be different ways of expressing the same idea. There is another way in which they may be related – perhaps there are some very general moral principles which underlie a number of more specific principles and concepts. 'Killing is wrong' is a moral principle, and thus, in accordance with our definition, it is a general rule which applies to specific cases. But there is a sense in which it is also specific – it tells us about a particular *class* of actions which are wrong, rather than providing a principle of such generality that we can work out from it whether *any* proposed action is wrong. Is there a general principle of this kind?

In this chapter we shall examine two moral theories each of which offers its own distinctive, and supposedly over-arching, moral principle. These are not the only moral theories which philosophers have put forward, but we concentrate on them because the principles they offer are clearly addressed to individuals, and intended as a basis for decisions as to what is the right thing to do in any particular situation. This chapter thus gives some insight into moral theory and also the opportunity to practise the skill of applying principles to specific cases. You may wish to increase your knowledge of moral theory by studying two other important theories – the social contract idea of John Rawls (J. Rawls, *A Theory of Justice*, Cambridge: Harvard University Press, 1971), which was mentioned in Chapter 5 in relation to the concept of rights, and the idea of 'virtue ethics', dating from Aristotle, which emphasises the importance of good character, and which is briefly discussed in James Rachels' *The Elements of Moral Philosophy*, McGraw-Hill, 1993. We shall not cover these theories in detail because 'virtue ethics' does not offer a single clear principle, and although Rawls theory includes principles of justice, these are intended as a basis for social organisation rather than a basis for individuals to make their ethical decisions.

In Chapter 2 we suggested the following strategy for assessing principles: think about some of the cases to which the principle must apply; consider whether any of these applications shows that there is something wrong with the principle; think about the way in which the principle should be modified. We shall take a practical approach of this kind in order to evaluate the two moral theories. We shall also consider the relationship of each theory to the moral concepts we have discussed – in particular to the concepts of harm, justice and autonomy, which we have suggested are crucial to reasoning well about practical ethical issues. We shall not discuss the way in which each of these theories seeks to *justify* its major principle. If you want this kind of detail about the theories, you can read J.S. Mill's *Utilitarianism* (1863) and H.J. Paton's *The Moral Law: Kant's Groundwork of the Metaphysic of Morals* (1948). For our purposes, it is sufficient to note that the principles proposed by these theories have some plausibility, and help to illuminate discussions in practical ethics.

Utilitarianism

The moral theory of utilitarianism originated with Jeremy Bentham in the late eighteenth century, and was advocated by John Stuart Mill in the nineteenth century, but this does not mean that it is out-dated and of merely historical interest. Many present-day philosophers apply the

framework of utilitarianism to practical ethical issues such as abortion, euthanasia, war and so on.

Clarifying the principle

The crucial feature of utilitarianism is its insistence that the rightness or wrongness of an action is determined by its consequences for everyone affected by it – the theory is described as a consequentialist theory. The early utilitarians claimed that the best consequences would be those which contained the greatest amount of happiness. John Stuart Mill's statement of the utilitarian principle was: 'Utility, or the Greatest Happiness Principle, holds that actions are right in proportion as they tend to promote happiness, wrong as they tend to produce the reverse of happiness' (J.S. Mill, *Utilitarianism*, 1863 – Chapter 2).

Modern utilitarians, instead of talking about 'happiness', are likely to say that what we should aim for is the maximum satisfaction of interests or of preferences. Peter Singer, for example, suggests that we should 'adopt the course of action most likely to maximise the interests of those affected' (Singer 1993: 13), and also that we should judge actions 'by the extent to which they accord with the preferences of any beings affected by the action or its consequences' (Singer 1993: 94).

The definitions raise a number of questions. First, is it the *actual* consequences (as Mill's definition suggests), or the *expected* consequences (which Singer implies) which determine the rightness of an action? If utilitarianism is to be a practical guide to conduct, then since we cannot in advance know the actual consequences of actions, it is most sensible to regard the principle as referring to expected consequences. Second, should we be judging what *a particular action* is likely to bring about (as Singer suggests) or should we (as Mill seems to think) judge whether *actions of this kind* generally bring about the best consequences? We shall return to this question later in this chapter when we discuss act-utilitarianism and rule-utilitarianism.

The definitions also illustrate a problem which arises from the classical version of the theory as expressed by John Stuart Mill, which tells us to maximise happiness. The problem is that it is not clear whether we should aim to give people what *we* think will make them happy – i.e. to give them what is judged to be in their interests – or to give them what *they* say will make them happy – i.e. to satisfy their expressed preferences. In 1994 in the Netherlands, where euthanasia is tolerated for terminally ill patients, a doctor was prosecuted for administering a lethal drug to a woman who was severely depressed after the deaths of her two sons and said she no longer wanted to live. The doctor acted to satisfy her preference, but this

may not have been in her interests if she could have been helped through her depression. However, Peter Singer clearly thinks that, in general, consideration of interests and consideration of preferences will produce the same result, since he talks about: 'the plausible move of taking a person's interests to be what, on balance and after reflection on all the relevant facts, a person prefers' (Singer 1993: 94).

Let us assume that we can generally safely judge interests on the basis of preferences, and think about how we would apply the utilitarian principle 'Act so as to produce maximum satisfaction of preferences of those affected by an action'. Is it now clear what we have to do? Obviously we have to think about how our actions will affect others as well as ourselves, and what those others will prefer. But who or what counts as 'others'? Presumably we must take into account anyone or anything capable of having preferences. So if animals have preferences, even if plants have preferences, we must take those preferences into account.

Applying the utilitarian principle

First, try for yourself to apply the utilitarian principle to the following examples. In each case, work out what the principle tells you is the right thing to do, and consider whether this shows that there is something wrong with the principle.

Example 1: Umberto, the plastic surgeon

> Umberto is a plastic surgeon specialising in the repair of birth defects. He lives in a part of the country where there is no-one else with his skills and qualifications. Umberto has two children, whom he sees very little of because he spends long hours at the hospital. He is a nice father and his children want to see more of him. He would not earn significantly less money if he decided not to work at weekends and to take holidays with his children. And his children would be happier. But hundreds of other children would then not get the operations that they need in order to live normal lives.
>
> (Adam Morton, 'Teaching Philosophy', *Cogito*, Spring 1994, p. 76)

Example 2: Marietta's uncle

> Marietta looks after her old uncle who has become mean and miserly, and wants to change his will. In his more generous days, he had bequeathed a vast sum of money to a charity which supports both cancer research and the provision of hospices for children with terminal cancer. This money would be of great benefit in helping to find a cure for childhood cancers, and in making the lives of those who cannot be cured much happier. But the uncle has decided to leave his money

instead to his estranged son who lives abroad, is already rich, and is much less generous than his father. The uncle is very forgetful about his medication, and often would take an overdose if Marietta did not stop him. He has arranged to see his solicitor to change his will. On the night before this meeting is due to take place, Marietta sees him about to take an additional dose of medication, which she knows will be fatal. If he dies in the night, the money will go to the charity, and the lives of many children will be made happier.

Example 3: Sam and famine relief

Having recently graduated with a degree in social studies, Sam has landed a very well paid job. She has no dependants, enjoys life very much, and would like to use her new found affluence in order to travel whenever possible, to buy a car, to go to the theatre often, and to eat at expensive restaurants. But she sees an appeal on television for donations to famine relief, and begins to think that she should forgo some of her pleasures in order to increase the happiness of people in other parts of the world. How much should she give? How many of the people who are suffering does she have to worry about. Does she have to think about the welfare of future generations, as well as the welfare of people alive now?

Umberto must maximise the satisfaction of the preferences of everyone affected by the amount of time he spends on his job. If he spends less time on his job, hundreds of children will not get what they would prefer – their birth defects remedied by Umberto's skills in plastic surgery. If he continues to spend the same amount of time, or even more time on his job, his own two children will not get what they would prefer – to spend more time with their father. More preferences will be satisfied if Umberto continues to work at weekends, so the utilitarian principle tells him that that is what he should do.

Is there a problem with this recommendation? Well, we can certainly identify something which would pull Umberto in the other direction, and this would not just be a *wish* to consider the interests of his own children, but a sense that he has a *moral obligation* to regard his own children's interests as a very high priority. The utilitarian principle does not allow him to give more weight to the interests of those closest to him. In Jeremy Bentham's words, 'Each person is to count for one, nobody for more than one'. So utilitarianism has difficulty allowing for special moral obligations associated with roles such as those of parent, brother, sister and so on.

What must Marietta do, according to the utilitarian principle? More preferences will be satisfied if she allows her uncle to take the overdose of medicine, thereby preventing him from changing his will in such a way that the cancer charity will not benefit from his wealth. But would this be the right thing to do? Wouldn't it be unfair to her uncle to let him take an accidental

overdose, if she would not have treated another elderly person, whose death could not have benefited others, in the same way? The problem for utilitarianism highlighted by this example is that the principle of maximising preferences can result in some people's preferences, even for such important things as freedom and continued existence, being denied in the interests of a larger group.

What does the utilitarian principle tell Sam that she must do about contributions to famine relief? She needs to find out how her contributions would be used, and the extent to which others would benefit. But certainly she must consider the satisfaction of the preferences of people all over the world, and it is likely that more preferences will be satisfied if she gives a large percentage of her income to charity, and does not satisfy her own preferences for a relatively luxurious life. Does she have to consider the preferences of future generations of people? Nothing in the principle rules this out. If more people will be born, and will have preferences, then these preferences will have to be taken into account.

The implications of the utilitarian principle begin to seem too demanding when we realise that it obliges us to take account of the entire world and future generations, even if we have not been wondering, like Sam, whether we should contribute to famine relief. Perhaps utilitarians can modify their principle in such a way as to restrict the preferences we should take into account to those of our immediate neighbourhood, or social group, or country. Or maybe the implication that we have very wide obligations is correct, and our tendency to regard the principle as too demanding is due to selfishness or limited sympathy. We leave you to think about this. However, whichever conclusion is correct, there is still a problem for utilitarianism in relation to our obligations to future generations. If one's own country, or the world, could provide a comfortable life for a larger population than at present, then it is likely that more preferences could be satisfied if the population were increased. This suggests that not only should we try to maximise satisfaction of preferences of those who happen to be born in the future, we should also try to maximise the number of people born who would then have preferences which we could satisfy. This would generate a moral requirement to produce children, rather than it being left to individuals to choose the size of their family. This rather odd implication could be avoided if the utilitarian principle were to read: 'Act so as to produce maximum satisfaction of preferences of those now alive who will be affected by your action'. However, this restriction would mean that we did not need to worry about the effects, for example, of pollution or of exhaustion of resources on future generations, and perhaps we should worry about such things.

Can the principle be modified in order to deal with the problems raised by Examples 1 and 2 – i.e. the difficulty of allowing for special responsibilities

such as those of parents, and the suggestion that utilitarianism cannot safeguard justice? The standard response by utilitarians is to say that reference to moral rules must be incorporated into the principle. Instead of recommending that we should simply perform the action which produces the best consequences, they recommend that we should follow the rules of conduct which, if they were followed by all, would generally produce the best consequences. This is a move from what is called act-utilitarianism to what is called rule-utilitarianism. The rule-utilitarian principle would be: 'Act in accordance with the rules of conduct which tend to produce maximum satisfaction of preferences'. (As we said earlier, this seems to be what Mill's principle recommends.) The rule-utilitarian will point out that most people will prefer to have rules of conduct which require parents to give greater priority to the interests of their own children, and which forbid treating some people badly in order to improve the fortunes of others. Most people will prefer a rule which safeguards justice, because they would feel insecure and fearful if there were no such rule.

Whilst avoiding some of the pitfalls of act-utilitarianism, rule-utilitarianism runs into its own problems. One criticism is that it isn't really utilitarian, because it has acknowledged that maximising the satisfaction of preferences might not be the most important consideration. In Marietta's case, for example, no-one will know if she fails to prevent her uncle's death, so this particular flouting of the rule will not lead to general insecurity. Yet rule-utilitarianism would require her to save her uncle, and thereby satisfy fewer preferences, so that it seems to be inconsistent with the basic idea of utilitarianism. This, of course, does not necessarily mean that it is a mistaken theory – simply that it is radically different from act-utilitarianism.

Another problem is that it seems to give the wrong sort of justification for treating people justly. For example, if Marietta should not let her uncle take an overdose, then it seems mistaken to say that the reason for this is that it is required by a rule which satisfies the preferences of most people. Shouldn't she treat him justly even if the preferences of others would not require this?

Utilitarianism and moral concepts

How does the utilitarian principle relate to some of the concepts which have come up in our examples of arguments and decisions in earlier chapters?

Utilitarianism is clearly related to the concept of harm, and the idea that there is something wrong with causing harm to others, though it goes beyond that to imply that we should make positive efforts to produce beneficial outcomes for others. It stresses the idea of beneficence – doing positive good.

Yet it may have to permit harm to some individuals in order to benefit a majority.

Its relationship with some of the other concepts we have considered, for example, justice, rights and autonomy is less direct. The concept of rights could be derived from rule-utilitarianism, as follows. We will probably find that most people are happier if there is general observance of a rule which implies that human beings have certain rights, some of which will be protected by law. So your having the moral right to life, or the moral right to liberty is justified by the fact (if there is such a fact) that preferences will be maximised if everyone is regarded as having these rights. The concepts of justice and of autonomy could be accommodated in a similar way by rule-utilitarianism.

The concept of a person is not of central importance to utilitarianism, because it does not restrict moral concern to persons. Nevertheless, just as those who use the concept of a person need some view about which entities have the capacity of self-consciousness, so also does utilitarianism depend on views about the *capacities* of various entities – in this case the capacity to have preferences. So if utilitarians are to be able to settle questions about how different entities should be treated, they need to be able to say which kinds of entities can have preferences.

Kant's moral theory

Immanuel Kant, a German philosopher writing in the late eighteenth century produced a moral theory firmly opposed to the idea that the morality of an action is determined by its consequences. Consequentialist moral theories, such as utilitarianism, are often contrasted with deontological moral theories, of which Kant's theory is an example, and which insist that some actions are right and some actions are wrong regardless of their consequences. The word 'deontological' derives from the Greek word meaning 'duty'.

Kant said that the way to find out what our duties are is to apply a principle which he called the categorical imperative – 'Act only according to that maxim by which you can at the same time will that it should become a universal law' (*Groundwork of the Metaphysic of Morals*, 1785). This connects with the idea that you shouldn't act in any way in which you would be unwilling to allow everyone else to act, which reminds us that if we think it is wrong for others to behave in a certain way, we cannot consistently claim that it is morally acceptable for us to behave in that way in exactly the same circumstances. You may remember that we made the same point in Chapter 3 about Saul's decision as to whether to become a vegetarian – that if he thought that the widespread practice of meat-eating was wrong on the grounds that it involved killing animals, he could not then excuse his own

meat-eating by saying 'It will make no difference to how many animals are killed if *I* give up eating meat'.

The wording of the principle may tempt us to think that Kant's theory depends on individual preferences, if we interpret 'can will' as meaning 'would be willing to allow'. This interpretation would mean, for example, that those who were willing to be treated cruelly would be morally justified in being cruel to others. However, Kant almost certainly meant 'can rationally and without inconsistency will', and he would have thought that someone who did not mind cruel treatment was irrational. There is, nevertheless, a difficulty in deriving from the categorical imperative the moral rules which Kant thought could be derived – for example, that murder, lying and breaking promises are wrong. Let us illustrate this in relation to murder. Someone contemplating murder must ask 'Can I rationally and without inconsistency will that others act in this way?' To insist that the answer must be 'no' requires an assumption that it is irrational both not to mind being murdered, and to be willing to take the risk that one may be murdered. These may be reasonable assumptions. However, the greatest problem arises when we think about how the action which the potential murderer contemplates must be described. Imagine a jealous husband who considers whether it is morally permissible to murder his wife's lover, and assume that the husband loves his wife so much that he would never become the lover of someone else's wife. Would it be irrational of him to will that all jealous husbands may murder their wives' lovers? Perhaps not, and if not then the test of the categorical imperative would imply that it would not be wrong to commit murder. But surely it would be wrong.

Kant offered another principle, called the formula of the end in itself, which he thought was simply another version of the categorical imperative. Yet, it is different, in that it makes a more specific recommendation about how we should treat other people. It says 'Act in such a way that you always treat humanity, whether in your own person or in the person of any other, never simply as a means but always at the same time as an end'. We shall concentrate on this principle, rather than the categorical imperative, because it is less abstract and thus easier to apply to examples of moral dilemmas. Indeed, the principle underlies some of the concepts used in discussions of practical ethical issues, but in order to see this we need a clear idea of what the principle means.

Clarifying the formula of the end in itself

The first thing to notice about the principle is that it tells us how we should treat 'humanity', so it has no implications for the treatment of animals. Kant himself thought that we have no direct duties to animals, and that animals

exist in order to serve the ends of humans. However, we can evaluate the principle as applied to human beings without accepting Kant's own views of the status of animals. Moreover, we should note that Kant suggested that we have *indirect* duties to animals, to the extent that cruelty to animals may make us callous towards human beings.

The principle gives us two requirements; we must never treat other people 'simply as a means', and we must always treat any other person 'as an end'. In the first of these, the word 'simply' is crucial, because in many transactions in everyday life, we use others as a means to our own ends. For example, we use the shop assistant as a means to making our purchases, we use the doctor as a means to restoring our health. What is meant by using someone *simply* as a means is trying to get them to do things, for our own purposes, which they would not choose to do if they were fully informed. We do not manipulate the shop assistant and the doctor if they are carrying out their jobs voluntarily. But it is possible to manipulate others by deceiving them, and this is one of the things which the principle tells us we must not do. Another example of using others simply as a means is suggested by the story of Marietta and her uncle. Suppose Marietta had been so determined that her uncle's money should go to the cancer charity that she had decided to kill him herself in order to be absolutely certain that he could not change his will. She would then have been guilty of simply using him as a means to benefit others.

The other part of the principle is more obscure. What does it mean to say we should treat someone 'as an end'? Modern Kantians interpret this as meaning that we should not merely respect others as rational persons with aims and purposes of their own, but that we should also make some attempt to help others to achieve some of those aims. This begins to sound like the utilitarian requirement to aim to satisfy the preferences of others, but it differs in that it does not demand that we *maximise* anything. In Onora O'Neill's words: 'Kantians will claim that they have done nothing wrong if none of their acts is unjust, and that their duty is complete if in addition their life plans have in the circumstances been reasonably beneficent' ('A Simplified Account of Kant's Ethics' in *Matters of Life and Death*, ed. T. Regan 1986).

Applying the formula of the end in itself

Before reading further, think about the three examples which we considered in relation to utilitarianism. What would the formula of the end in itself tell Umberto, Marietta and Sam that they should do?

In Umberto's case, the formula offers little to help him make a decision. We can assume that since he is a nice father, and since his work gives

children the help which they and their parents want, he is using neither his own children, nor his patients simply as a means to his own ends. His dilemma is that if he continues to work long hours, he will not promote his children's aim to spend more time with their father, and if he works shorter hours he will not promote the aim of some other children to get treatment. The principle itself cannot give an answer as to whose aims should take priority.

We have already suggested that Marietta would be acting against the formula if she killed her uncle as a means to making the lives of child cancer sufferers happier. Would she be using him as a means if she did not actively kill him, but just neglected to stop him taking an overdose? That question may be difficult to answer, but it does seem that she would be failing to treat him as an end, i.e. as a person with aims of his own. We assume that he wants to go on living, and it is clear that he wants his wealth to go to his son. By failing to stop him taking an overdose, she would be failing to help him achieve these aims.

What does the formula tell Sam that she should do? She knows that there are people dying from famine. She is not doing anything to them, so she cannot be said to be using them as a means to her ends. But if she does not contribute to famine relief, is she failing to treat these people as ends? The question is difficult to answer partly because Sam never comes into contact with the people who are starving. How can we be treating people in any way at all, or even failing to treat them in a proper way, if we never meet them? Does Kant's theory have nothing to say about ethical issues concerning people with whom we have no direct contact?

In 'Kantian Approaches to Some Famine Problems', Onora O'Neill outlines a way in which Kant's formula can be seen as relevant to such issues. It contrasts with Peter Singer's utilitarian position, and in order to get a detailed understanding of this contrast, you could read O'Neill's paper and Chapter 8 in Singer's *Practical Ethics*. In that chapter, Singer sets out an argument as follows:

> **First premise: If we can prevent something bad without sacrificing anything of comparable significance, we ought to do it.**
>
> **Second premise: Absolute poverty is bad.**
>
> **Third premise: There is some absolute poverty we can prevent without sacrificing anything of comparable moral significance.**
>
> **Conclusion: We ought to prevent some absolute poverty.**

You can practise your skills of argument assessment on this example. He

assumes that the contributions of individuals would make a difference to those in other countries suffering absolute poverty, and he suggests 'that those earning average or above average incomes in affluent societies, unless they have an unusually large number of dependants or other special needs, ought to give a tenth of their income to reducing absolute poverty'.

O'Neill's discussion is not so directly addressed to individuals, but she does talk about duties of beneficence in times of famine. Her view is that, because, according to Kant's theory, the most basic duty of beneficence is to put people in a position to be autonomous, and because no-one who is starving can be autonomous, the relief of famine rates high among the duties of beneficence. It is not clear whether this duty is owed principally by governments rather than individuals.

In a final section of her paper, entitled 'Lifeboat Earth', O'Neill draws what looks like a stronger conclusion. She points out that the economic interdependence of countries is now such that the lifestyles and activities of Western countries have effects on the well-being of other countries. So affluent countries cannot simply claim that in giving but little to famine relief they are meeting all the obligations that Kantian theory demands, because if our lifestyles actually cause starvation in other countries, then we are killing people in other countries (a matter of failing to be just) rather than merely failing to save them, and thereby failing to be as kind to them as we could be. This makes famine relief more an issue of justice than of beneficence.

Kant's theory and moral concepts

The two requirements of the formula of the end in itself connect directly with two of the crucial concepts for practical ethics, namely justice and autonomy.

Our simple definition of justice in Chapter 4 was 'treating equals equally', and we have suggested that utilitarianism has difficulty in accommodating this value, since in some situations maximum satisfaction of preferences might be achieved by treating an individual or a minority group differently from the majority. Kant's formula requires that everyone be treated with equal respect as persons with aims and purposes of their own. The idea of rights does not appear in Kant's theory. His focus was on duty – what we owe to others, rather than what we are entitled to demand from others. There is, of course, some relationship between rights and duties, as we pointed out in Chapter 5. If you have a right to life, then others have a duty not to kill you. But the relationship may not work the other way round. So it is not obvious that the concept of rights can be derived from Kant's formula.

The concept of autonomy underlies Kant's theory. He thought that human beings are all potentially rational and capable of making choices about their own lives. The formula of the end in itself implies that autonomy should be respected – we should not manipulate others into doing things they would not otherwise have

wished to do. O'Neill's interpretation of 'treating others as ends' also implies that autonomy should be promoted – we should make some attempt to put some others in a position to be autonomous. Kant did not explicitly introduce the concept of a 'person', but his emphasis on the importance of autonomy implies that our moral obligations are owed only to those entities which are potentially rational and autonomous.

There is a connection between Kant's theory and the concept of harm. The formula requires that we treat others as persons with aims of their own, and if we can assume that it would be contrary to anyone's aims that they should be harmed, it does suggest that we should not directly inflict harm on other people.

Yet Kant's theory does not require us to take into account all the potentially harmful consequences of our actions, since it insists that some actions are right irrespective of their consequences. One example can illuminate this. The formula generates the moral rule that we must tell the truth, since if we deceive others, we are not respecting their autonomy. In Chapter 2 we used the example of someone being asked the whereabouts of Anne Frank's family when they were hiding from the Nazis in Amsterdam. The consequences of telling the truth in this situation could include suffering and death for the members of the Frank family. But Kant's insistence that the moral rules which are derived from his principle are absolute implies that we must not lie, even if refusing to lie results in the infliction by others of terrible harm. Kant defends this stance by reminding us that we cannot be certain what the consequences of our actions will be – perhaps a lie will not have the good consequences which we predict. As we pointed out in our discussion of famine relief, Kant's theory does include the concept of beneficence – doing good to others, but it is not specific about how much good we should do and to whom.

Moral theories and some other principles

The principle that 'Killing is wrong'

The principle that killing is wrong underlies many of the arguments about ethical issues, for example abortion, euthanasia and vegetarianism, which we have encountered in earlier chapters. It also figures in discussions about the morality of capital punishment and of war. As a moral consideration, it has something in common with the concept of harm, in that it is very widely accepted, and we would think there was some moral defect in someone who simply could not see that it was a moral consideration at all – someone who said 'What's so bad about killing?' It is tempting to think that we do not even have to seek a justification for the principle. However, because the principle is so important, any moral theory which offers a supposedly over-arching moral principle should be able to justify the idea that killing is wrong by means of its own principle.

We suggested in Chapter 2 that you might attempt to assess the principle 'Killing is wrong' by thinking about its applications. You may have concluded that some exceptions need to be built into the principle. For example, it may not be wrong to kill someone who is threatening the life of another person, if that would be the only way to save the potential victim. Should this idea of the permissibility of killing in self-defence be extended to make killing in wartime morally acceptable? Could there be other exceptions to the principle, to allow, for example, abortion and euthanasia to be morally permissible? What kind of justification for the principle, and what exceptions, are implicit in the two moral theories we have examined?

Utilitarianism and 'Killing is wrong'

Utilitarians must justify the claim that killing is wrong in terms of the consequences of killing – principally that it deprives someone of a future during which their interests and preferences could have been satisfied. Act-utilitarianism cannot insist on an absolute rule against killing, because in some circumstances the satisfaction of preferences may be maximised if some people are killed. Rule-utilitarians would justify an absolute rule against killing by pointing out that if there were no such rule, everyone would feel very insecure and fearful. Yet, there could be specific exceptions, even to an absolute rule.

Kant's moral theory and 'Killing is wrong'

Kant's theory can justify the claim that killing is wrong by means of the concept of respect for persons and for their autonomy.

To kill someone who wanted to go on living would be to fail to treat that person as someone with aims and purposes of their own.

We have said that Kant believed that moral rules were absolute, so he would not have accepted that the consequences of killing could ever justify it. However, there is room for the idea of an exception to the principle in some cases. For example, killing in self-defence could be seen as treating others in the way in which they have decided people may be treated. Interestingly, Kant uses this idea to justify capital punishment.

The doctrine of acts and omissions

In this section we ask you to consider whether some controversial cases should be exceptions to the principle that killing is wrong. Read each of the following examples, and decide whether it would be wrong to kill in these cases. If so, try to say exactly why.

Example 1: Permanent vegetative state

A young man, Tony Bland, suffered brain damage due to oxygen starvation when he was crushed against the barrier at Hillsborough football ground in April 1989. He was judged to be in a permanent vegetative state – his brain stem was alive, so that he could breathe normally, but all his higher brain function had gone, so that he could not feed himself, and had to be connected to a feeding tube in order to survive. After almost four years, his parents applied for a legal judgement to allow him to be disconnected from the feeding tube, and therefore to die. This was granted, the tube was disconnected, and he died about ten days later. At the time, his doctor said that he could see no moral difference between giving Tony a lethal injection and removing the feeding tube, but he would not do the former because it was illegal.

Example 2: Elderly Alzheimer sufferers with pneumonia

Many elderly people in nursing homes who suffer from Alzheimer's disease are not treated with antibiotics when they fall ill with pneumonia. They could be treated, and often they would recover, but a judgement is made that it is appropriate that their lives should not be prolonged in this way. The pneumonia is allowed to take its course, and death results. Would it be wrong to give a lethal injection in these circumstances?

Example 3: Infants with severe disabilities

Imagine the case of a child being born with a disability which is so severe that the child will die within a month, and there is no medical procedure which can remedy the disability. Would it be wrong to kill the child?

In the first two cases, a decision is effectively made that a life should come to an end, and in the third case the end of life is not preventable. You may have thought that in each case, it is right to let the person die, but it would not be right to kill them. Those who think that killing someone is always morally worse than letting someone die are relying on a general claim, known as the **acts and omissions doctrine**, which says that:

> There is a moral difference between *performing an act* which has certain consequences, and *failing to act* when that failure to act has exactly the same consequences.

You can probably see immediately that utilitarians must reject this doctrine, because it is the consequence which determines the rightness of our behaviour, so doing nothing can be just as bad as doing something. This is why utilitarianism imposes such a heavy responsibility upon us to relieve famine. Kant's theory is harder to apply to omissions than to actions, since his focus was upon how we know whether *intentional* actions are right or wrong. This was why we had difficulty deciding whether those who do not contribute to famine relief can be said to be 'treating' the starving in any way at all.

The acts and omissions doctrine seems plausible because we can think of many examples in which an omission *does* seem less bad than an action. For example, it seems morally worse to drown someone by pushing them into a river than to fail to rescue someone who has fallen into the river; and it seems morally worse to send poisoned food to people who are starving than to fail to send them food at all. But should we accept it as a general rule or principle applying to all cases? We leave you to consider this in the light of your responses to our examples.

However, the Bland case prompts an important observation about the acts and omissions doctrine – that it can be difficult to know whether someone's behaviour should be described as an act or as an omission. Tony Bland's feeding tube had to be disconnected from time to time in order to clean it. Once the decision had been made that he should be allowed to die, the tube could have been removed immediately (clearly an action). But in fact the doctor waited until the tube was next due to be disconnected, and he did not re-connect it. Was this an omission, rather than an action? And what moral difference could this distinction make?

The doctrine of double effect

We now focus on some cases of deaths which result from a particular action, rather than an omission, but which are claimed not to be cases of intentional killing, even though the person performing the action knows that death will result. Think about whether the person is doing something wrong.

Example 4: Painkilling injections

People dying from cancer are often in such severe pain that they require very large doses of pain-killing drugs. Such drugs are extremely toxic, and can bring about death earlier than the illness would have done. This practice was referred to by Melanie Phillips in her article presented in Exercise 4, where she said that doctors 'administer pain relief which

might have the side-effect of hastening the death of an already dying patient', and she describes this as 'absolutely in line with a doctor's commitments to preserve life and relieve suffering'. Her comments imply that doctors acting thus are doing nothing wrong even though they know that death will result from their actions.

Example 5: Killing 'human shields'

When Saddam Hussein thinks that Western nations are likely to attack Iraq's weapons installations, he arranges for the presence of 'human shields' at these sites. Civilians will occupy the buildings in order to deter enemies from bombing them. In times of war, would it be wrong to bomb these installations in order to destroy Iraq's weapons, even though it was known that this would also kill civilians?

In the first example, the doctor kills the patient as a side-effect of an action which is aimed at relieving the patient's pain; in the second example the bomber kills the civilians as a side-effect of an action which is aimed at destroying the enemy's weapons. Should we regard these cases as instances of killing which are exceptions to the principle that killing is wrong?

The claim that there can be exceptions in such circumstances is often defended by what is called the **doctrine of double effect**, which says:

One need not be held responsible for those effects of one's actions, which, though foreseen, are not intended, provided that:

(i) the action performed is done because it will have some good effect, even though it may also have bad effects, and

(ii) one intends only the good effects and not the bad effects of the action, and

(iii) the bad effect is not the means by which the good effect is achieved. (An example of what would be ruled out by condition (iii) would be Marietta killing her uncle as a means to the good effect of benefiting cancer sufferers. In order to justify an action by means of the doctrine, the bad effect really must be a side-effect.)

Utilitarianism has no need of this doctrine, because it judges the rightness or wrongness of actions on the basis of their consequences, and not the intentions behind them. The doctrine is likely to be used by non-consequentialists as a way of discounting what seem to be strong moral objections to certain courses of action. Should we accept the doctrine as applying to all cases in which there are bad side-effects?